Summary

Exit

Question 1:
Your company operates nationally and plans to use GCP for multiple batch workloads, including some that are not time-critical. You also need to use GCP services that are HIPAA-certified and manage service costs. How should you design to meet Google best practices?

A. Provisioning preemptible VMs to reduce cost. Discontinue use of all GCP services and APIs that are not HIPAA-compliant.
B. Provisioning preemptible VMs to reduce cost. Disable and then discontinue use of all GCP and APIs that are not HIPAA-compliant.
C. Provision standard VMs in the same region to reduce cost. Discontinue use of all GCP services and APIs that are not HIPAA-compliant.
D. Provision standard VMs to the same region to reduce cost. Disable and then discontinue use of all GCP services and APIs that are not HIPAA-compliant.

Type : SINGLE SELECTION

Answered By
1
Correct Answers Given By
1
Percentage
100 %

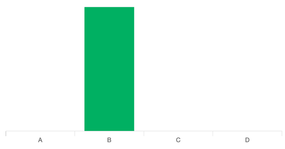

Question 2:
You have an App Engine application that needs to be updated. You want to test the update with production traffic before replacing the current application version. What should you do?

A. Deploy the update using the Instance Group Updater to create a partial rollout, which allows for canary testing.
B. Deploy the update as a new version in the App Engine application, and split traffic between the new and current versions.
C. Deploy the update in a new VPC, and use Google's global HTTP load balancing to split traffic between the update and current applications.
D. Deploy the update as a new App Engine application, and use Google's global HTTP load balancing to split traffic between the new and current applications.

Type : SINGLE SELECTION

Answered By
1
Correct Answers Given By
1
Percentage
100 %

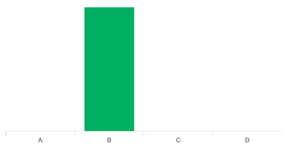

Question 3:
You are migrating third-party applications from optimized on-premises virtual machines to Google Cloud. You are unsure about the optimum CPU and memory options. The application have a consistent usage patterns across multiple weeks. You want to optimize resource usage for the lowest cost. What should you do?

A. Create a Compute engine instance with CPU and Memory options similar to your application's current on-premises virtual machine. Install the cloud monitoring agent, and deploy the third party application. Run a load with normal traffic levels on third party application and follow the Rightsizing Recommendations in the Cloud Console.
B. Create an App Engine flexible environment, and deploy the third party application using a Docker file and a custom runtime. Set CPU and memory options similar to your application's current on-premises virtual machine in the app.yaml file.
C. Create an instance template with the smallest available machine type, and use an image of the third party application taken from the current on-premises virtual machine. Create a managed instance group that uses average CPU to autoscale the number of instances in the group. Modify the average CPU utilization threshold to optimize the number of instances running.
D. Create multiple Compute Engine instances with varying CPU and memory options. Install the cloud monitoring agent and deploy the third-party application on each of them. Run a load test with high traffic levels on the application and use the results to determine the optimal settings.

Type : SINGLE SELECTION

Answered By
1
Correct Answers Given By
1
Percentage
100 %

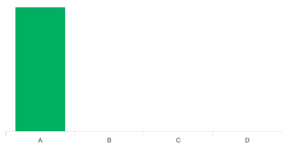

Question 4:
Your company is running its application workloads on Compute Engine. The applications have been deployed in production, acceptance, and development environments. The production environment is business-critical and is used 24/7, while the acceptance and development environments are only critical during office hours. Your CFO has asked you to optimize these environments to achieve cost savings during idle times. What should you do?

A. Create a shell script that uses the gcloud command to change the machine type of the development and acceptance instances to a smaller machine type outside of office hours. Schedule the shell script on one of the production instances to automate the task.
B. Use Cloud Scheduler to trigger a Cloud Function that will stop the development and acceptance environments after office hours and start them just before office hours.
C. Deploy the development and acceptance applications on a managed instance group and enable autoscaling.
D. Use regular Compute Engine instances for the production environment, and use preemptible VMs for the acceptance and development environments.

Type : SINGLE SELECTION

Answered By
1
Correct Answers Given By
1
Percentage
100 %

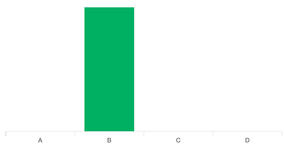

Question 5:
You have created several pre-emptible Linux virtual machine instances using Google Compute Engine. You want to properly shut down your application before the virtual machines are preempted.What should you do?

A. Create a shutdown script named k99.shutdown in the /etc/rc.6.d/ directory
B. Create a shutdown script registered as a xinetd service in Linux and configure a Stackdriver endpoint check to call the service
C. Create a shutdown script and use it as the value for a new metadata entry with the key shutdown-script in the Cloud Platform Console when you create the new virtual machine instance
D. Create a shutdown script, registered as a xinetd service in Linux, and use the gcloud compute instances add-metadata command to specify the service URL as the value for a new metadata entry with the key shutdown-script-url

Type : SINGLE SELECTION

Answered By
1
Correct Answers Given By
1
Percentage
100 %

Question 6:
You are analyzing and defining business processes to support your startup's trial usage of GCP, and you don't yet know what consumer demand for your product will be. Your manager requires you to minimize GCP service costs and adhere to Google best practices. What should you do?

A. Utilize free tier and sustained use discounts. Provision a staff position for service cost management.
B. Utilize free tier and sustained use discounts. Provide training to the team about service cost management.
C. Utilize free tier and committed use discounts. Provision a staff position for service cost management.
D. Utilize free tier and committed use discounts. Provide training to the team about service cost management.

Type : SINGLE SELECTION

Answered By
1
Correct Answers Given By
1
Percentage
100 %

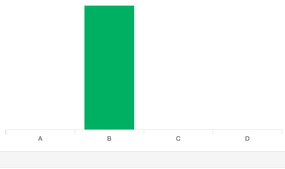

Question 7:
You have an application that waits for images to be uploaded into storage. Whenever an image arrives, a thumbnail of the image is created. The thumbnail can be displayed on a web page. Users can see the full image when they click on the thumbnail. The code is written in JavaScript and uses Node.js. What service might be the easiest and cheapest way for deploying this program?

A. Compute Engine
B. App Engine Standard
C. GKE and Kubernetes
D. Cloud Functions

Type : SINGLE SELECTION

Answered By
1
Correct Answers Given By
1
Percentage
100 %

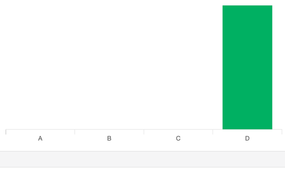

Question 8:
Your company has an aging data center that is running short of capacity. It wants to move applications to the cloud quickly to avoid refreshing the hardware. There are many applications to be moved, most are written in Java or C#. Most applications are currently running in virtual machines using VMware. What strategy would you recommend they follow when moving to Google Cloud?

A. Use a lift and shift approach to get applications running in Compute Engine virtual machines quickly. Then, selectively rewrite applications to take greater advantage of cloud services.
B. Prioritize the applications. Then, rewrite them in priority order to run in App Engine taking greatest advantage of cloud services.
C. Get everything running in Docker containers first. Then, use Kubernetes and Google Container Engine for deploying applications.
D. Only use the cloud for new development. Leave old apps where they are as moving them into the cloud is too prone to errors. Retire old apps in the current data center to get more capacity.

Type : SINGLE SELECTION

Answered By
1
Correct Answers Given By
1
Percentage
100 %